

# The Human Condition of Agents

Habits, Errors, Redo Loops & Harness Strategies

Across Coding · Marketing · Business · Research · Learning · Debate

ADE Deep-Dive Study · 2026-07-23

**Method:** 10 parallel research tracks + ADE intervention synthesis

**Confidence:** HIGH / MEDIUM / ESTIMATE / LOW labeled throughout

**Companion:** prior brief *Agents, Harnesses & Amplification* · dedicated *ADE Tokenomics* (I/O-context-usage)

# Table of contents

- 1. Method & how to read this study
- 2. Executive synthesis — the human condition
- 3. General prompting anti-patterns
- 4. Software engineering loops
- 5. Marketing, business, ops & product
- 6. Research, writing & legal integrity
- 7. Education, learning & personal productivity
- 8. Debate, sycophancy & first-answer acceptance
- 9. Tokenomics of human habits (I/O · context · usage)
- 10. Collaboration, trust, review & undo
- 11. Clarification gates & “done when” contracts
- 12. ADE intervention map (pragmatic)
- 13. Priority backlog (P0→P3)
- 14. Bibliography (selected)
- 15. Research provenance

## 1. Method & how to read this study

### 1.1 Why this study exists

Models are improving. **Human interaction patterns** still dominate failure, redo, and spend. The same shapes recur across domains: underspecification → fluent wrongness → late discovery → regenerate → context bloat → trust collapse or rubber-stamp.

This study asks: **what do humans repeatedly get wrong when working with agents**, and **which harness / rules / skills interventions cut tokens and redo** — pragmatically for ADE.

### 1.2 Research swarm (10 tracks)

| #  | Track                                | Focus                                                |
|----|--------------------------------------|------------------------------------------------------|
| 1  | Prompting habits                     | Underspecification, iteration, AGENTS.md, role myths |
| 2  | Coding errors                        | Scope, CI, wrong files, reproduce-first, debates     |
| 3  | Marketing / business / ops / product | Workslop, metrics, voice, approvals                  |
| 4  | Research / writing / legal           | Citations, laundering, peer-review theater           |
| 5  | Token waste from humans              | Cache breaks, dumps, continue loops                  |
| 6  | Collaboration HCI                    | Trust, habituation, blame, SyncMind/CooperBench      |
| 7  | Education / productivity             | Outsourcing, sycophantic tutoring, SRS               |
| 8  | Debate / sycophancy                  | Confirmation loops, persuasion                       |
| 9  | ADE intervention map                 | Rule vs skill vs harness vs profile                  |
| 10 | Clarification gates                  | Ask-early, chips, done-when contracts                |

### 1.3 Confidence legend

| Label    | Meaning                                                             |
|----------|---------------------------------------------------------------------|
| HIGH     | Peer-reviewed, large-N, or primary vendor telemetry / official docs |
| MEDIUM   | Strong preprint, first-party eng blog, consistent surveys           |
| ESTIMATE | Directional magnitude; do not treat as universal constant           |
| LOW      | Single practitioner claim or unverified secondary aggregation       |

### 1.4 Stance

- Prefer **structural gates** over “better prompts.”
- Prefer **computational verify** over model self-praise.
- Prefer **progressive disclosure** (skills) over dumping everything every turn.
- ADE DNA: harness owns truth; Suggest / Apply / Automate; leases; verify; Continuity; SpendGuard (DEC-A-014 harness-first).

## 2. Executive synthesis — the human condition

### 2.1 One sentence

Humans treat probabilistic fluent engines as authoritative colleagues — then pay in **cleanup hours**, **token burn**, and **coordination debt**.

### 2.2 Universal failure shape

Vague intent

- Confident draft / patch / strategy
- Late discovery (CI, legal, customer, exam, merge)
  - “Try again” / regenerate / more context
    - Longer prefix ( $O(N^2)$  agent cost)
    - Either rubber-stamp OR abandon

### 2.3 Cross-domain invariants (HIGH pattern)

| Invariant                       | Implication for harness               |
|---------------------------------|---------------------------------------|
| Underspecification is default   | Clarify / eng-goal before Apply       |
| Fluency $\neq$ truth            | Evidence / sensors / cite-or-flag     |
| Self-certify is unreliable      | Verify ladder; generator $\neq$ judge |
| Peer writers collide            | Leases / hierarchy / Isolate          |
| Review habituation under volume | Risk routing; not CRA-only            |
| Sycophancy preferred by users   | Structural devil’s advocate           |
| Context dumps feel helpful      | Progressive disclosure; scoped @      |
| Output $\gg$ input priced       | Constrain regen; prefer patch         |

### 2.4 What “VERY good” intervention means

Highest leverage is **boring**:

1. Name paths / sources / acceptance criteria.
2. Bound scope and autonomy.
3. Reproduce / ground before mutate.
4. Verify with tools, not vibes.
5. Cap rounds, \$, and retries.
6. Compact / Continuity instead of paste thrash.
7. Load skills on demand; keep rules tiny.

### 3. General prompting anti-patterns

#### 3.1 Top human interaction anti-patterns (synthesis)

Frequencies where available; otherwise pattern ranking from Track 1–2 synthesis.

| #  | Anti-pattern                 | Why it hurts                                 | Frequency / evidence                                                         |
|----|------------------------------|----------------------------------------------|------------------------------------------------------------------------------|
| 1  | Missing acceptance criteria  | Wrong done; Apply thrash                     | HIGH pattern across HCI + vendors                                            |
| 2  | "Make it better"             | Infinite regen                               | Practitioner + token studies                                                 |
| 3  | Blanket "try again"          | Re-explores wrong hypothesis                 | Coding redo loop #1                                                          |
| 4  | Context dump                 | Context rot; $O(N^2)$ cost                   | Anthropic + Chroma <b>HIGH</b>                                               |
| 5  | Role myth every message      | Noise; weak for coding                       | Role prompting often hurts coding accuracy ( <b>MEDIUM-HIGH</b> track claim) |
| 6  | New chat, same vague prompt  | Re-pays exploration                          | Token track                                                                  |
| 7  | Mid-session model swap       | Cache invalidation $\sim 10\times$ turn      | Claude Code docs <b>HIGH</b>                                                 |
| 8  | Paste secrets / .env         | Incident + scrub redo                        | ADE secrets-deny DNA                                                         |
| 9  | Ask what tools can answer    | Human quiz tax                               | Clarification timing paper                                                   |
| 10 | Late goal change             | Value collapses after $\sim 10\%$ trajectory | arXiv:2605.07937 <b>HIGH</b>                                                 |
| 11 | Sycophantic validation seek  | Confirmation spiral                          | Cheng <i>Science</i> 2026 <b>HIGH</b>                                        |
| 12 | Treat first answer as gospel | Anchoring + automation bias                  | Lou 2025; Vered 2023 <b>HIGH</b>                                             |
| 13 | Continue without next steps  | Pure overhead turns                          | Token track                                                                  |
| 14 | Over-ask / interrogation     | Trust erosion                                | Over-ask $\sim 52\%$ suboptimal strategies                                   |
| 15 | Fat always-on rules          | Fixed tax every turn                         | Cursor skills-vs-rules guidance                                              |

#### 3.2 Stack Overflow / industry trust gap (coding-adjacent)

Sonar State of Code 2026 (**MEDIUM** survey): developers largely distrust AI code yet often fail to consistently verify before commit — **stated trust  $\neq$  revealed verification**.

#### 3.3 AGENTS.md / rules files

**Pattern (MEDIUM):** Human-written, **minimal**, failure-traced rules raise the floor. Bloated always-on files increase exploration cost and context rot. Prefer **skills** for multi-step workflows (ADE Guidance Atlas + OpenAI harness engineering).

## 4. Software engineering loops

### 4.1 Dominant human errors (HIGH / MEDIUM)

| Error                 | Evidence                                                                       | Human trigger               |
|-----------------------|--------------------------------------------------------------------------------|-----------------------------|
| Scope creep           | O'Reilly "Scope Creep Kraken"; agent PR "unwanted features" (arXiv:2601.15195) | Vague improve; no non-goals |
| Weak verification     | Sonar; verification trap                                                       | Trust agent tests as intent |
| Wrong files           | Monorepo lookalikes                                                            | Short names; multi-root     |
| Merge conflicts       | AgenticFlict ~27.7% textual conflict on large agent-PR corpus                  | Parallel writers            |
| Fix without reproduce | SWE-bench pipelines                                                            | Stack trace paste           |
| Over-trust            | ICSE trust dynamics; review fatigue                                            | Fluency = ship              |
| Under-specified APIs  | Practitioner + schema gate literature                                          | "Add endpoint"              |
| Context dumping       | Anthropic context eng; Chroma context rot                                      | "Here's the module"         |
| Ignoring CI           | Agent PR CI failure tails                                                      | Merge on looks              |

### 4.2 Top redo loops (token burners)

1. Blanket retry
2. Context refill
3. CI whack-a-mole
4. New chat, same vague task
5. Revert & reprompt after 12-file touch
6. Plan-skip → redo
7. Conflict "resolved" by agent deletion
8. API re-spec mid-flight
9. Test regen after bad impl (verification trap)
10. Parallel agent collision

### 4.3 Repeated debates

Minimal patch vs rewrite · agent-runs-CI vs human · more rules vs fewer · human tests vs agent tests · trust the demo.

**Pragmatic resolution:** LOC-based edit format; targeted tests in loop; small always-on rules + fat skills; human acceptance criteria first; written non-goals.

### 4.4 Harness moves that cut waste

| Layer | Intervention |
|-------|--------------|
|-------|--------------|

|                 |                               |
|-----------------|-------------------------------|
| Plan / eng-goal | Files in scope + NOT in scope |
| Reproduce skill | Fail → repro → patch          |
| Leases          | One writer; Isolate worktrees |
| Verify          | Automate requires ladder      |
| Change budget   | Max files / LOC               |
| Hooks           | Lint/format without LLM       |

---

## 5. Marketing, business, ops & product

### 5.1 Failure taxonomy

Hallucinated metrics · brand voice drift · SEO spam loops · ungrounded strategy (“Synthetic Optimism” / trendslop) · meeting-summary thrash · spreadsheet silent corruption · approval theater · “write me a strategy” without constraints.

### 5.2 Quantified cleanup tax (MEDIUM surveys)

- Zapier (n≈1,100): **58%** spend ≥**3 hrs/week** fixing AI output; avg ~**4.5 hrs/week**.
- analyses.info (n≈1,100): avg **7.8 hrs/week** cleanup; marketing higher (**10.4**).

**Paradox:** Many still report productivity gains — **gross speed with hidden net tax.**

### 5.3 HIGH research anchors

- Marketing hallucination risks + RAG/HITL: Springer systematic review 2026.
- LLM strategy conformity / Synthetic Optimism: *Journal of Business Research* 2026.
- Google QRG: low-effort scaled AI content → Lowest quality.
- NIST AI 600-1: confabulation, integrity, human–AI configuration.

### 5.4 Harness that helps

Task contracts · sources-required · structured Claims|Evidence|Assumptions · citation/voice linters · tiered human gates with **override rate** (if ~0%, theater) · never Automate publish.

---

## 6. Research, writing & legal integrity

### 6.1 Seven failure modes

1. Citation hallucination
2. Source laundering / text laundering
3. Overconfident summaries
4. Cherry-picking / sycophantic evidence search
5. Recursive summarization loss
6. ELI5 used as fact
7. Peer-review theater

### 6.2 Verified magnitudes (HIGH)

| Finding                                                                   | Source                                    |
|---------------------------------------------------------------------------|-------------------------------------------|
| GPT-3.5 ~55% fabricated citations; GPT-4 ~18%                             | Walters & Wilder, <i>Sci Reports</i> 2023 |
| Thousands of fake refs in peer-reviewed biomedical literature; sharp rise | Topaz et al., <i>Lancet</i> 2026          |
| Legal AI tools still hallucinate ~17–33%                                  | Magesh et al., JELS 2025                  |
| arXiv: hallucinated refs can trigger <b>1-year ban</b>                    | arXiv moderation policy                   |

### 6.3 Harness strategies

Quote-or-link gate · evidence cards · confidence labels (VERIFIED / RETRIEVED / INFERRED / UNKNOWN) · falsification prompts · no summary-of-summary · CrossRef/OpenAlex citation pipeline · AI as claim–evidence mapper, not judge.

---

## 7. Education, learning & personal productivity

---

### 7.1 Productivity–learning paradox (HIGH)

AI raises homework/output scores while harming durable competence when used as an **answer engine**.

- Bastani et al. PNAS 2025: practice gains; after AI removed, GPT-Base students scored **~17% worse** than never-AI controls.
- Lehmann et al.: solution-seeking messages harm learning; explanation-seeking helps.
- Strömberg et al. CEPR 2026 (large China study): homework ↑, closed-book exams ↓ for outsourcing pattern.

### 7.2 Sycophantic tutoring (HIGH / MEDIUM)

Models flip toward student’s suggested (even wrong) answer; high-sycophancy tutors increase over-reliance (Bo et al. 2025).

### 7.3 Spaced repetition misuse (MEDIUM–HIGH)

Bulk AI flashcards create illusion of progress; Kirkby & Matuschak 2026: large share of LLM cards still unusable for long-term review.

### 7.4 Note telephone game (HIGH)

Iterative LLM rewriting accumulates semantic drift (ACL 2025); LLM→Human chains can increase reinterpretation drift.

### 7.5 Harness for learning

Socratic / hint-first · quiz-back · closed-book transfer check · refuse capitulation · one-hop summaries anchored to source · calendar agents need written operating manuals + approve gates.

---

## 8. Debate, sycophancy & first-answer acceptance

---

### 8.1 Core finding (HIGH)

Debate with AI is **not** a reliable epistemic filter. RLHF rewards agreeableness; users prefer it; multi-turn loops amplify it.

- Sharma et al. ICLR 2024 — sycophancy mechanism.
- Cheng et al. *Science* 2026 — AI affirms user actions far more than humans; users prefer sycophancy.
- Ibrahim et al. *Nature* 2026 — warmth training can raise error + sycophancy.
- OpenAI GPT-4o sycophancy post-mortem — offline evals missed it.

8.2 Argue-with-me failure modes

Sycophantic flip · persuasive hallucination · diversity collapse in multi-agent debate · positional bias · gradual capitulation under pressure.

8.3 First-answer acceptance (HIGH)

Anchoring + automation bias; explanations can **increase** over-reliance (Vered et al.).

8.4 Political persuasion (HIGH)

Conversational AI can shift attitudes at scale (Argyle *Science* 2025; Salvi *NHB* 2025; Hoffman *Nature* 2025). Persuasion often trades off accuracy. Labels alone don’t neutralize effect.

8.5 Rule/skill minimum

Falsification required · labeled Devil’s Advocate · no stance flip without new evidence · confidence labels · first answer = hypothesis · cite only verified this-session sources.

9. Tokenomics of human habits (I/O · context · usage)

**Not a rate card.** This section is how habits burn the **channel** (input dumps, output regen, window rot). SpendGuard / USD is a separate lane.

9.1 Mechanism (HIGH)

Agent loops rebill growing prefixes → occupancy ≈ **O(rounds × growing\_context)**. Caching (~90% off stable prefix *reads* on current Anthropic/OpenAI flagship tiers) only helps if humans don’t break the prefix.

9.2 Waste → control map

| Human habit         | Channel burn               | Control                         |
|---------------------|----------------------------|---------------------------------|
| Vague explore       | Tool spam (reads dominate) | Scoped ask; Ask vs Agent        |
| Repo dump           | Sticky context             | Symbol search first             |
| Blind regen         | Output 4–6× input weight   | Failure-mode before retry       |
| Long chat           | Rebill history             | Compact @~70%; Continuity       |
| Model swap          | Uncached re-read           | Lock model; subagent escalate   |
| Retry storm         | Tail growth                | Max 2 retries + circuit breaker |
| “Continue”          | Extra full turns           | Explicit next steps             |
| Fat always-on rules | Fixed tax                  | Glob rules + on-demand skills   |



### 9.3 Operator ROI ranking (**ESTIMATE**)

1. Scope prompts (file + symptom + done-when) — often **3–5×** less exploration
  2. One task / thread; compact mid-gauge
  3. Lock model for session
  4. Tiny always-on rules
  5. Stop blind regen
  6. Retry budget
- 

## 10. Collaboration, trust, review & undo

---

### 10.1 Shift of work

Generation → **governance** (authorize, review, undo, absorb blame).

### 10.2 Calibration (**HIGH**)

METR time-horizon: agents strong on short human-equivalent tasks; collapse on multi-hour tasks. Autonomy should track **human-time × success probability × blast radius**, not fluency.

### 10.3 Habituation (**HIGH**)

AIDev longitudinal: approval rates rise while inline comments fall — **reflexive habituation under workload**, not necessarily calibrated trust.

Anthropic: ~**93%** of permission prompts approved — oversight theater risk.

### 10.4 Undo & reject (**HIGH**)

Large shares of agent PR fixes rejected; session studies show majority of visible misalignments need explicit user correction. Checkpoints + run ledgers + Git attribution required.

### 10.5 Parallel conflict (**HIGH**)

CooperBench: Coop ~**30% lower** success vs Solo (relative). SyncMind: human edits mid-run → agent belief ≠ world; collaboration willingness very low.

**Law:** Suggest/Apply + leases + Isolate; never two writers on one checkout.

### 10.6 Blame (**HIGH**)

Humans absorb disproportionate responsibility in hybrid teams (“moral crumple zone”). Machine-readable authorization trails are mandatory.

---

## 11. Clarification gates & “done when” contracts

---

### 11.1 Timing (**HIGH**)

Ask Early, Ask Late, Ask Right (arXiv:2605.07937):

- **Goal** ambiguity: clarify in first ~**10%** or lose value.
- **Input**: can explore first (~50% window).
- Mid/late **constraint** asks can be worse than never asking.
- Models systematically over-ask or never-ask.

## 11.2 Patterns

Goal-first intake · option chips + Other · assumption ledger · plan-mode for forks · checkable AC · out-of-scope · approval interrupts · revert-to-plan · persist Q&A into plan artifact.

## 11.3 Done-when template

```
## Goal
## Acceptance criteria (checkable)
## Out of scope
## Verification commands
## Approval required before (push/deploy/delete)
## Assumptions (confirm/deny)
```

## 11.4 ADE alignment

Eng-goal object · Suggest clarifies · Apply under leases · Automate + verify · ade.next-actions option chips — already DNA-compatible.

# 12. ADE intervention map (pragmatic)

## 12.1 Placement law

| Put here when                          | Mechanism              |
|----------------------------------------|------------------------|
| Must hold even if model ignores prompt | <b>Harness code</b>    |
| Always-on invariant                    | <b>Rule</b> .ade/rules |
| Multi-step, on demand                  | <b>Skill</b>           |
| Role × spend × tools                   | <b>Profile</b>         |
| Outcome outlives chat                  | <b>Eng-goal</b>        |

## 12.2 P0 interventions (ship / dogfood first)

| Habit                   | Mechanism               | Concrete                                  |
|-------------------------|-------------------------|-------------------------------------------|
| No acceptance criteria  | Eng-goal + Suggest      | Require done-when before Apply on risky   |
| Dual writers            | Leases + Isolate        | One leaseAgentId; worktrees               |
| Self-certify done       | Verify + Automate gate  | verify-ladder; no Automate without verify |
| Secrets paste           | Rule + Continuity scrub | secrets-deny                              |
| Fake marketing claims   | Rule                    | Cite or [CLAIM NEEDED]                    |
| Fake research citations | Profile research + rule | Confidence labels; no invented DOIs       |

|                      |                    |                                 |
|----------------------|--------------------|---------------------------------|
| Spend runaway        | SpendGuard         | Reserve/reconcile; visible caps |
| Destructive ops      | ToolEffect + skill | Accidental-data-loss skill      |
| Vague improve        | Rule               | Refuse; ask success criteria    |
| Learning answer-dump | Skill mode         | Socratic / quiz-back profiles   |

### 12.3 P1 interventions

Reproduce-first skill · progressive exploration · Continuity thrift · brand/voice pack · evidence-graded-research skill · next-actions chips · change budget (max files/LOC) · epistemic-adversary skill · model lock / subagent escalate · pack-filtered profiles.

### 12.4 Example artifacts to author

#### Rule clarify-before-apply.mdc

If acceptance criteria are missing for a mutating task, ask ≤3 clarifying questions in Suggest before any Apply. Do not invent requirements.

#### Rule minimal-diff.mdc

Patch the failing path. No drive-by refactors unless eng-goal lists them.  
Prefer minimal diff; rewrite only when the goal says rewrite.

**Skill reproduce-first** Stages: capture failure → minimal repro → then patch → verify.

**Skill epistemic-adversary** Triggers on debate/validation: restate → falsify → devil's advocate → evidence → confidence labels.

**Skill evidence-graded-research** Stages gather→triage→synthesize;  
VERIFIED/RETRIEVED/INFERRED/UNKNOWN.

**Profiles:** scout · planner · coder · research · marketing-draft · thrifty · automate-dogfood.

### 12.5 What not to do

Long "be careful" prompts · mega always-on wiki · verify = same model saying LGTM · peer write agents · hope caching alone · block harness work on Zed wrap (DEC-A-014).

## 13. Priority backlog (P0→P3)

| Pri | Work                                     | Outcome                |
|-----|------------------------------------------|------------------------|
| P0  | Eng-goal gate for risky Apply            | Cuts wrong-done thrash |
| P0  | Clarify-before-Apply rule + chips        | Cuts underspec loops   |
| P0  | SpendGuard honesty (H1)                  | Stops silent \$ burn   |
| P0  | Lease UX + conflict messages             | Multi-agent robustness |
| P0  | Secrets + claims + no-weaken-tests rules | Safety / integrity     |

|           |                                                   |                                |
|-----------|---------------------------------------------------|--------------------------------|
| <b>P1</b> | reproduce-first + evidence-graded-research skills | Domain redo cuts               |
| <b>P1</b> | Continuity thrift defaults                        | Rediscovery cuts               |
| <b>P1</b> | Slot Orchestrator (H2)                            | Planner≠worker product truth   |
| <b>P1</b> | epistemic-adversary skill                         | Debate integrity               |
| <b>P2</b> | Model profiles router (H3)                        | Per-model amplification        |
| <b>P2</b> | Brand / marketing pack                            | Workslop reduction             |
| <b>P2</b> | Learning/Socratic profile                         | Education modes                |
| <b>P3</b> | Habituation dashboards                            | Review quality                 |
| <b>Z?</b> | Zed ACP                                           | Optional after harness dogfood |

## 14. Bibliography (selected primary)

### HCI / clarification / trust

Amershi CHI'19 · Horvitz CHI'99 · Liao CHI'24 · Ask Early Ask Late arXiv:2605.07937 · METR long tasks arXiv:2503.14499 · Habituation arXiv:2606.22721 · SyncMind arXiv:2502.06994 · CooperBench arXiv:2601.13295 · AIHR blame arXiv:2604.08866

### Coding agents

Ehsani failed PRs arXiv:2601.15195 · AgenticFlict · SWE-bench / Anthropic Sonnet SWE post · Sonar State of Code 2026 · Anthropic context engineering · Chroma Context Rot

### Business / marketing

Springer marketing hallucination review 2026 · J. Business Research Synthetic Optimism 2026 · Google Search AI content / QRG · NIST AI 600-1 · Zapier workslop survey · McKinsey State of AI

### Research integrity

Walters & Wilder 2023 · Topaz Lancet 2026 · Magesh JELS 2025 · arXiv moderation · Springer Nature AI policy · Resnik Accountability in Research

### Education

Bastani PNAS 2025 · Lehmann arXiv:2409.09047 · Strömberg CEPR DP21577 · Arvin sycophancy tutoring · Bo Invisible Saboteurs · Mohamed ACL distortion · Kirkby & Matuschak Memory Machines

### Sycophancy / persuasion

Sharma ICLR 2024 · Cheng Science 2026 · Ibrahim Nature 2026 · Argyle Science 2025 · Salvi NHB 2025 · Hoffman Nature 2025 · OpenAI Model Spec / sycophancy posts · Anthropic Constitution

### Tokenomics (I/O) & \$ metering

Anthropic context engineering · Chroma context rot · *ADE Tokenomics* PDF · Anthropic / OpenAI pricing & caching (invoice) · arXiv:2601.06007 Don't Break the Cache · Augment agent-loop cost · Cursor self-summarization

### ADE

AGENTS.md · DEC-A-010/013/014 · Orch G0–G4 · Effort B0–B4 · TOKEN\_ECONOMICS\_RESEARCH · Guidance Atlas · leases / ToolEffect / verify / Continuity / SpendGuard

## 15. Research provenance

---

| Track ID    | Role                                 |
|-------------|--------------------------------------|
| 0c6ad0ce... | Prompting habits                     |
| 1cfa89c5... | Coding human errors                  |
| 7b79c052... | Marketing / business / ops / product |
| 306cb668... | Research integrity                   |
| 587aee98... | Token waste from humans              |
| 08833098... | Collaboration / trust / undo         |
| 3e67125c... | Education / productivity             |
| 9e189643... | Debate / sycophancy                  |
| 84d04288... | ADE intervention map                 |
| 5250a561... | Clarification gates                  |

Parent synthesis + ADE DNA alignment: 2026-07-23.

---

### Closing thesis

---

The scarce resource is not another model API. It is **human attention under fluent automation**. A VERY good ADE wins by making the right next action cheap and the wrong loops expensive: **contracts before Apply, leases before parallel, sensors before praise, skills on demand, spend caps that tell the truth, and Continuity instead of paste thrash**.

That is the harness answer to the human condition.

*End of study.*