

ADE Master Gameplan

How to build ADE from the research stack

Harness OS · Tokenomics · Compaction · Contracts · Slots · Honest \$ · Optional eyes

Edition: 2026-07-23 · **Schema:** ade.master-gameplan/v1

Canvas: ADE-master-gameplan.canvas.tsx

Locked by: DEC-A-010 / 013 / 014 · AGENTS.md DNA

Inputs: Ideal ADE · What ADE Could Do · Tokenomics · Compaction/Fertility · Market study · Harness brief

0. One sentence

Build ADE as a **local harness OS** that makes the research wins unavoidable — contracts before Apply, envelopes for side effects, leases for writers, verify outside the generator, honest spend, Continuity + compaction for the token channel — with Desktop as control plane and Zed as **optional eyes**, never identity.

1. North star (non-negotiable)

Principle	Source
Quality $\approx$ model $\times$ harness $\times$ human process	Harness brief / market study
Product identity = harness , not IDE fork	DEC-A-014 · AGENTS.md
Suggest $\approx$ Propose · Apply $\approx$ Act under leases · Automate $\approx$ Act + verify	DNA
Planner $\neq$ Worker $\neq$ Verifier · no peer democracy	CooperBench · DEC-A-014
Tokenomics $\neq$ \$ · Effort $\neq$ SpendGuard $\neq$ window	Tokenomics paper
Compaction = harness job · private “ADE language” = avoid	Compaction deep-dive
Side effects consume authorized envelopes	Ideal ADE

Equation for builders:

Ship discipline as product objects – not as prompt pleading.

2. Research → product map

Research finding	Product object	Track
Underspec → redo	Eng-goal + ≤ 3 clarify before Apply	G1
Dual writers / peer tax	Slot Orchestrator + leases + claim-one	H2
Self-certify	Verify ladder + Automate gate	Have → deepen
Context dump / $O(N \times \text{rounds})$	Tool-blob clearing + capsules @ boundary	C1–C2
Fertility $\neq$ invent cipher	Catalog IDs · diffs · compact JSON	C / Tokenomics
Model roulette	ade.model-profile/v1 + router	H3
Opaque \$	SpendGuard honesty	H1
Review habituation	Risk-tiered HITL	G2
Continuity > paste	Capsules + thrift + host next_safe	Have → C3
Editor chrome temptation	Soft shell after H+C dogfood	Z?


3. Baseline (already shipped — do not rebuild)

Area	Status
------	--------

Orch G0–G4	Suggest / Apply / Isolate / eng-goals / Role strip
Effort B0–B4	Rounds + output caps · honesty · Continuity Continue
Leases + worktrees + task claim	On-disk truth
Verify G0–G5 + Automate verify-on-complete	Shipped
Rules / skills / pack profiles · T0–T3 activate	Shipped
Continuity ade.handoff/v1	Shipped (thrift resume)
SpendGuard reserve/reconcile (basic)	Partial honesty
Desktop Ideal spine	Control plane
ACP / Zed	Probe-only scaffold
Gold ~51	Deterministic; needs races

Stance: Amplify spine. Don’t reboot as an editor project.

4. Four parallel tracks (one critical path)



Critical path (DEC-A-014 + research ROI):
G1 contract gate || H1 SpendGuard honesty || C1 tool-blob clearing
 Then **H2 slots → H3 profiles → C2/C3 compaction → H4/H5 → Z?**

Pick **one** primary eng focus per sprint; keep the other two as thin parallel slices if staffing allows.

5. Wave plan

Wave 0 — Lock (done)

- Accept DEC-A-014
- Research stack published (Ideal / Could Do / Tokenomics / Compaction / Market)
- Dogfood Continuity + Isolate scripts green enough to build on

Wave 1 — Force multipliers (~2–4 weeks)

Kill the highest-cost failure modes from human/market studies.

ID	Ship	Done when
----	------	-----------

G1	Contract before Apply	Act tools blocked without active eng-goal (AC+OOS+verify) or ≤ 3 clarify / logged waive
H1	SpendGuard honesty	No silent \$0 rates; Trust shows used/reserved/remaining $\approx$ invoice class; sane reserve
C1	Tool-result clearing	Stale blobs stubbed/cleared by policy; re-fetchable; gold proves mask $\geq$ baseline

Dogfood story: “I cannot Apply into the void; I cannot burn \$ I don’t see; long turns don’t drown in tool paste.”

Wave 2 — Amplification (~3–5 weeks)

ID	Ship	Done when
H2	Slot Orchestrator truth	Suggest=Planner · Apply=Worker · Verify=Judge bound to claim + heartbeats; dual-writer impossible
H3	Model profiles + router	ade.model-profile/v1; router $\neq$ None; “why this model” chip; no silent mid-task swap
C2	Boundary + @~70% capsules	Compact at task/verify boundaries; threshold backstop; structured sections
C3	Write-before-compact	Decisions/failing/paths land in .ade/ before summary
G2	Risk-tiered HITL	Secrets/infra/migrate/publish force Propose or confirm even under Automate; waivers logged

Dogfood story: “A slot crew with honest models and a clean window.”

Wave 3 — Robustness (~2–4 weeks)

ID	Ship	Done when
H4	Lease conflict UX	Blocked Apply → wait / Isolate / rotate CTAs
H5	Harness gold races	Dual-writer, wrong-slot, budget, spend-cap, Isolate merge in gold
C4	SelfCompact-style tool	compact_context + rubric (fire on resolve; suppress mid-derivation)
C5	Compaction/fertility gold	Mask vs summary vs capsule fidelity metrics
E1	Action envelopes in trace	effect-paths-lease-autonomy-goal id per grant → turn + Continuity


Dogfood story: “Production survival under parallel workers and long missions.”


Wave 4 — Optional eyes (after Wave 2 dogfood)

ID	Ship	Done when
Z1	ACP soft shell	Full stdio loop; Suggest/Apply/leases/verify mapped; Open-in-Zed
Z2	Fork ladder review	Promote L3/L4 only with written chrome gaps ACP+Desktop cannot close

Non-goal until then: Mission Control chrome, SSO/SCIM, VSCode revival, private ADE cipher language.

6. Architecture target (build toward)





Eyes change. Brain does not.

7. Channel strategy (Tokenomics + Compaction) — build rules

- 1. **Thin always-on · catalog · activate** — already partially shipped; tighten always_apply bloat.
- 2. **Mask before summarize** — Complexity Trap.
- 3. **Structured capsules** — intent · decisions · paths · failing · next · verify.
- 4. **Diffs / hunks / compact JSON** — not full dumps; not private stego language.
- 5. **Effort counts output**; provider usage bills \$; chars÷4 is assembly-only.
- 6. **Compact at semantic boundaries**; ~70% safety net.
- 7. **English + saturated formats** for operator I/O (fertility≠quality; don't translate to "save tokens").

8. Success metrics (ship when these move)

Metric	Direction
Apply without contract (or waive)	→ 0
Automate complete without verify	→ 0
Dual-writer lease violations	→ 0
Spend UI vs invoice delta	→ 0
Continuity resume without paste	↑
Peak tokens / successful Apply (C1–C2)	↓
Compaction artifact retention (paths/failing)	↑
Gold green incl. race cases	required
Mid-task silent model swap	→ 0

9. 90-day sequencing (recommended)

Days 1–14 (Wave 1 start)

Primary: **G1** contract gate. Parallel thin: **C1** clearing design + **H1** rate tables.

Days 15–35

Primary: **H1** honesty ship. Parallel: **C1** implement + gold. G1 dogfood harden.

Days 36–55

Primary: **H2** slot truth. Start **H3** schema. **C2** boundary capsules behind flag.

Days 56–75

Primary: **H3** router live. **C3** write-before-compact. **G2** risk HITL.

Days 76–90

H4 conflict UX · **H5** race gold · **C4** SelfCompact tool · decide Z1 go/no-go.

Adjust only with written ADR if Zed suddenly blocks dogfood (unlikely under DEC-A-014).

10. Team / work splitting (even if solo)

Hat	Owns
Gates	G1–G2, eng-goal UX, clarify chips
Orchestrator	H2, H4, leases/claim heartbeats
Money	H1 SpendGuard
Router	H3 profiles
Channel	C1–C5, Continuity polish
Eval	H5 + C5 gold
Eyes	Z1 only after Wave 2

Solo: sequence G1 → H1 → C1 → H2 → H3 → C2...

11. Anti-goals (explicit)

- Cursor/Zed clone as identity
 - Mission Control before slot truth
 - Flat peer agent swarms
 - Blocking harness on ACP
 - Private ADE “machineese” for tokenization
 - Caps that show \$0 while tokens burn
 - Automate that greenlights on model self-report
 - Dumping full chat into Continuity
-

12. Decision log hooks

If...	Then...
Invoice trust blocks overnight runs	Prioritize H1 over H2
Dual-writer dogfood fails weekly	Prioritize H2 over H3
Context rot kills long Apply	Prioritize C1/C2
Underspec redo dominates	Prioritize G1 (default)
Coding eyes block harness dogfood	Open Z1 early (exception)
Someone proposes ADE stego language	Point to compaction paper §5 — reject

13. Artifact index

Artifact	Role
This PDF	Master build plan
ADE-master-gameplan.canvas.tsx	Interactive board
Ideal ADE / What ADE Could Do	Normative + capability
Tokenomics / Compaction deep-dives	Channel science
Agents Meet Humans	Market habits
DEC-A-014	Priority law
Prior ADE-harness-multiagent-gameplan	H-track detail

14. Closing

ADE is already a harness spine with empty seats for honesty, gates, slots, and channel hygiene. The research says those seats — not editor chrome — are how you become the product that **survives production**.

Build order: contracts · honest \$ · clear the window · slot truth · profiles · compact well · then eyes.

End · *ade.master-gameplan/v1* · 2026-07-23